

AgPV Index

Exploratory data analysis for creation of an AG PV index: a composite scoring metric to identify USA counties with high cobenefits potential for agrivoltaics

- Data inputs: solar supply, weather hazards, energy burden, minority owned cropland

Data Processing

Things to do:

- aggregate NREL's solar supply data by county
- get relevant weather hazards data from csv file (hail and drought are positives, tornado is negative.)
 - might start ith overall agriculture burden for positives first
- load energy burden data
- get minoirty owned data from R2R indices to start

```
In [1]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
import os
```

Solar Supply Data NREL

Commonly cited NREL data for solar supply. I'm aggregating this by county.

<https://www.nrel.gov/gis/solar-resource-maps.html>

```
In [2]: # got solar supply data
solar_dir = 'solar-pv-reference-access-2023'
solar_file = 'reference_access_2030_moderate_supply-curve.csv'

solar_df = pd.read_csv(os.path.join(solar_dir, solar_file), dtype={'cnty_fips':str})
solar_df['cnty_fips'] = solar_df['cnty_fips'].apply(lambda x: x.zfill(5))

# get total for each county
location_cols = ['cnty_fips', 'county', 'state']
solar_sums_df = solar_df.groupby(location_cols, as_index=False).sum()

# I think we just need capacity_mw_ac
solar_sums_df = solar_sums_df[location_cols + ['capacity_mw_ac', 'capacity_mw_dc']]

# rename fips column
solar_sums_df = solar_sums_df.rename(columns={'cnty_fips':'FIPS'})
```

```
solar_sums_df.head()
```

```
Out[2]:
```

	FIPS	county	state	capacity_mw_ac	capacity_mw_dc
0	01001	Autauga	Alabama	22349.202978	29947.935193
1	01003	Baldwin	Alabama	46948.973155	62911.632997
2	01005	Barbour	Alabama	43344.657028	58081.848135
3	01007	Bibb	Alabama	28368.690972	38014.050777
4	01009	Blount	Alabama	18483.385380	24767.738851

NRI Weather hazards data

FEMA has created weather hazard index scores for various weather hazards. Right now, I'm only focusing on how these hazards affect agricultural losses. For AgPV, tornadoes are negative, since they destroy solar panels. Drought, hail, and heatwave are positives, since AgPV can help protect crops or diversify farmer income hagaintst these hazards.

Relevant columns:

- TRND_ALRA: expected tornado ag loss rate
- HWAV_ALRA: heatwave expected ag loss rate
- HAIL_ALRA: hail expected ag loss rate
- DRGT_ALRA: drought expected loss
- RESL_SCORE: community resilience score
- SOVI_SCORE: soviol vulnerability score

can use different suffixes for different metrics. Other metrics of interest:

- RISKV: risk value
- RISKS: risk index score
- RISKR: risk index rating

<https://hazards.fema.gov/nri>

```
In [3]: data_dir = 'AgPV_data'
nri_dir = 'NRI_Table_Counties'
nri_file = 'NRI_Table_Counties.csv'

nri_df = pd.read_csv(os.path.join(data_dir, nri_dir, nri_file), dtype={'STCOFIPS':s

# keep relevant columns
prefixes = ['TRND', 'HWAV', 'HAIL', 'DRGT']
suffixes = ['ALRA', 'RISKV', 'RISKS']
nri_keep_cols = ['_'.join([pf, sf]) for pf in prefixes for sf in suffixes]

nri_df = nri_df[['STCOFIPS', 'RESL_SCORE', 'SOVI_SCORE'] + nri_keep_cols]
```

```
# rename relevant columns
nri_df = nri_df.rename(columns={'STCOFIPS': 'FIPS'})
nri_df.head()
```

```
Out[3]:
```

	FIPS	RESL_SCORE	SOVI_SCORE	TRND_ALRA	TRND_RISKV	TRND_RISKS	HWAV_ALRA
0	01001	51.810001	51.299999	0.000029	2.831745e+06	73.846643	2.747418e-06
1	01003	86.120003	31.030001	0.000004	7.982720e+06	91.377665	9.080529e-07
2	01005	6.240000	99.269997	0.000009	2.845056e+06	73.942094	2.747418e-06
3	01007	19.730000	80.779999	0.000067	2.609927e+06	71.969456	4.200027e-06
4	01009	22.820000	51.369999	0.000077	6.930682e+06	89.723194	4.200027e-06

Energy Burden

Percent of household income spent on energy. AgPV can help high burden counties lower their energy burden.

<https://www.energy.gov/scep/slsc/lead-tool>

```
In [4]: eburden_file = 'LEAD Tool Data Counties.csv'

eburden_df = pd.read_csv(os.path.join(data_dir, eburden_file), skiprows=range(0,8),
eburden_df = eburden_df.rename(columns={'Geography ID': 'FIPS'})
eburden_df = eburden_df[['FIPS', 'Energy Burden (% income)']]
eburden_df.head()
```

```
Out[4]:
```

	FIPS	Energy Burden (% income)
0	01001	3
1	01003	2
2	01005	4
3	01007	4
4	01009	3

R2R Data

Farmer income and % of minority owned cropland analyzed for Roads to Removal. Don't get too sad at the minority owned farm numbers.

```
In [5]: farm_income_file = 'avg_farm_income.csv'

# 'Avg Farm Net Income ($)' has a bunch of wierd str values. '(D)' is undisclosed,
farm_income_df = pd.read_csv(os.path.join(data_dir, farm_income_file), dtype={'Coun
farm_income_df['FIPS'] = farm_income_df.apply(lambda x: x['State ANSI'].zfill(2) +
farm_income_df = farm_income_df[['FIPS', 'Avg Farm Net Income ($)']])
```

```
farm_income_df.head()
```

Out[5]:

	FIPS	Avg Farm Net Income (\$)
0	01001	18,279
1	01011	71,850
2	01047	35,071
3	01051	9,847
4	01063	19,870

In [6]:

```
# minority owned cropland
minority_crop_file = 'minority_owned_cropland_counties.csv'

min_crop_df = pd.read_csv(os.path.join(data_dir, minority_crop_file), dtype={'fips_
min_crop_df = min_crop_df.rename(columns={'fips_code': 'FIPS'})
min_crop_df = min_crop_df[['FIPS', 'Percent Minority Owned']]
min_crop_df.head()
```

Out[6]:

	FIPS	Percent Minority Owned
0	01005	0.000000
1	01011	1.984635
2	01013	0.000000
3	01025	0.000000
4	01043	0.415887

Merge all DFs together

In [7]:

```
all_dfs = [solar_sums_df, nri_df, eburden_df, farm_income_df, min_crop_df]
all_dfs = [df.set_index('FIPS') for df in all_dfs]
merged_df = pd.concat(all_dfs, axis=1)
merged_df = merged_df.dropna(subset=['county'])

# percent minority owned has lots of missing data (due to a lack of minority owned
merged_df['Percent Minority Owned'] = merged_df['Percent Minority Owned'].fillna(0)

# clean farm income data. This column is the messiest
def clean_farm_income(x):
    if isinstance(x, float):
        return x

    if x.strip() == '(D)':
        return np.nan

    else:
        return float(x.replace(',', ''))
```

```
def fill_na_max(ser: pd.Series):  
    return ser.fillna(ser.max())  
  
merged_df['Avg Farm Net Income ($)'] = merged_df['Avg Farm Net Income ($)'].apply(c  
  
# fill undefined rows (nan entries) with max value. This is most conservative estim  
merged_df['Avg Farm Net Income ($)'] = fill_na_max(merged_df['Avg Farm Net Income (  
  
# identify counties with negative net income  
merged_df['Negative Net Farm Income'] = merged_df['Avg Farm Net Income ($)'].apply(  
  
merged_df
```

Out[7]:

	county	state	capacity_mw_ac	capacity_mw_dc	RESL_SCORE	SOVI_SCORE
FIPS						
01001	Autauga	Alabama	22349.202978	29947.935193	51.810001	51.299999
01003	Baldwin	Alabama	46948.973155	62911.632997	86.120003	31.030001
01005	Barbour	Alabama	43344.657028	58081.848135	6.240000	99.269997
01007	Bibb	Alabama	28368.690972	38014.050777	19.730000	80.779999
01009	Blount	Alabama	18483.385380	24767.738851	22.820000	51.369999
...
56037	Sweetwater	Wyoming	401041.398994	537395.558058	30.709999	37.400002
56039	Teton	Wyoming	4915.661374	6586.986832	39.529999	19.190001
56041	Uinta	Wyoming	83816.558073	112314.202527	27.820000	40.639999
56043	Washakie	Wyoming	72277.627225	96852.034274	62.029999	26.610001
56045	Weston	Wyoming	155044.608219	207759.812413	9.740000	23.200001

3082 rows × 22 columns

Data Exploration


Are there strong correlations in the underlying variables? That might skew results.

Ac and DC solar potential are nearly identical and strongly correlated, so we can just use one of them.

Other than that, I didn't see any strong correlations between the variables. There is some positive correlation (~0.5) between tornado risk and the other weather hazards.


```
In [8]: # only use one set of nri metric to make the graph comprehensible
sns.set_theme(style='darkgrid')
plot_cols = [col for col in merged_df.columns if '_RISKV' not in col and '_ALRA' not in col]
sns.pairplot(merged_df[plot_cols])

# save plot
graph_dir = 'AgPV_graphs'
graph_file = 'agpv_pairplot.png'
plt.savefig(os.path.join(graph_dir, graph_file))
```



```
In [9]: fig, ax = plt.subplots(figsize=(20,20))
corr_df = merged_df.iloc[:, 2:].corr()
sns.heatmap(corr_df, annot=True, cmap='vlag', vmin=-1, vmax=1)
plt.title('AgPV Index Variables Correlation', size=20)

# save file
corr_file = 'agpv_corr.png'
plt.savefig(os.path.join(graph_dir, corr_file))
```



Index Preprocessing

NRI score values seem appropriate over using the value or expected loss metric since SVI and community resiliency also use score.

For roads to removal, I used a Box-Cox transform to try and fit to a normal distribution. This time, I'm trying quantile normalization, which really forces the data to a normal distribution. The results look much more gaussian. In a sense, it's mangling the data, but since we're already doing so much to it anyway with the scaling, it's probably OK. In most cases, we get very pretty bell curves.

Process:

- Transform underlying variables to fit normal dist'n

- Minmax scale data [0-1]
- Invert 'negative' variables (tornadoes and farmer income)
- Average values of underlying variables
- Minmax scal results [0-1]

```
In [10]: # use quantile normalization to make everything have a normal distribution
from sklearn.preprocessing import quantile_transform, MinMaxScaler
X = quantile_transform(merged_df[plot_cols].iloc[:, 2:], output_distribution='norma

# scale all values 0-1
scaler = MinMaxScaler()
X = scaler.fit_transform(X)

# create df from transformed data
normalized_df = pd.DataFrame(X, columns=merged_df[plot_cols].iloc[:, 2:].columns, i
normalized_df = normalized_df.drop(columns=['capacity_mw_dc', 'Negative Net Farm In

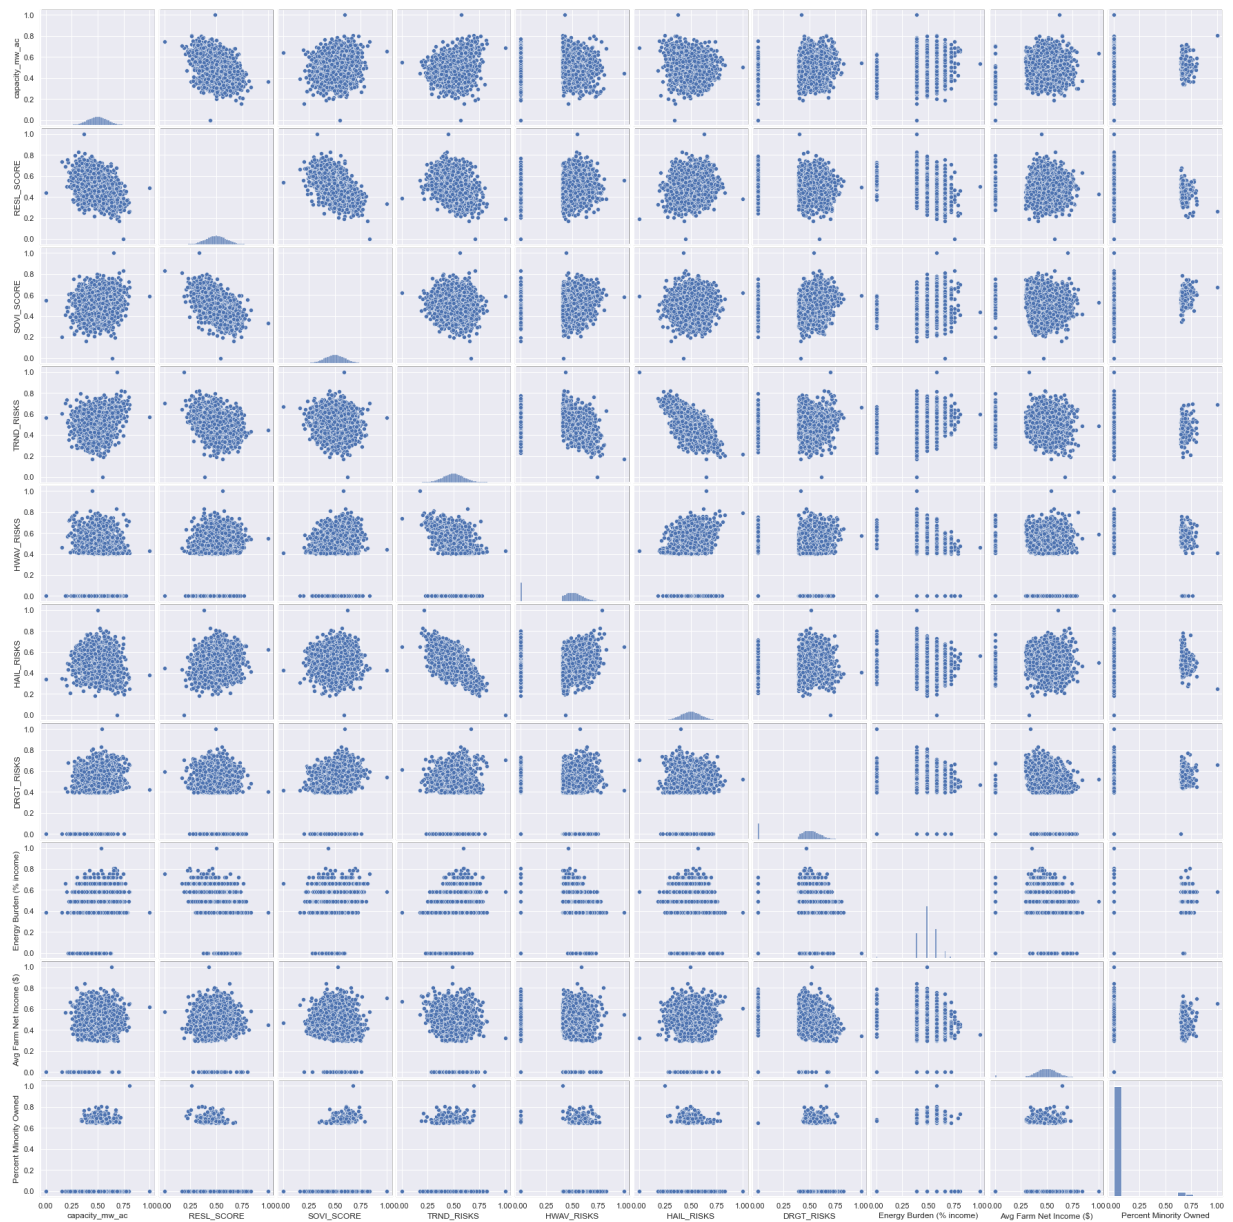
# need to invert the neagtives in the index: tornado risk and farm income
normalized_df['Avg Farm Net Income ($)'] = 1 - normalized_df['Avg Farm Net Income (
normalized_df['TRND_RISKS'] = 1 - normalized_df['TRND_RISKS']
normalized_df.head()
```

Out[10]:

	capacity_mw_ac	RESL_SCORE	SOVI_SCORE	TRND_RISKS	HWAV_RISKS	HAIL_RISKS
FIPS						
01001	0.547606	0.504096	0.503589	0.440039	0.607810	0.515519
01003	0.612906	0.603788	0.453189	0.369846	0.628607	0.505205
01005	0.606443	0.352327	0.735681	0.439753	0.548137	0.473010
01007	0.570571	0.418034	0.584009	0.445549	0.541421	0.482337
01009	0.530352	0.428265	0.503724	0.379363	0.588573	0.482416

```
In [11]: sns.pairplot(normalized_df)
```


Out[11]: <seaborn.axisgrid.PairGrid at 0x28b963424c0>



The transformations seemed to slightly exaggerate correlation between toronado and hail risk. Tornado and energy burden also have some correlation (0.3). This could cause some bias in the overall index.

```
In [12]: # correlation map with simplified variables and transformed
fig, ax = plt.subplots(figsize=(20,20))
corr_df = normalized_df.iloc[:, 2:].corr()
sns.heatmap(corr_df, annot=True, cmap='vlag', vmin=-1, vmax=1)
plt.title('AgPV Index Variables Quantile Transform Correlation', size=20)
```

```
Out[12]: Text(0.5, 1.0, 'AgPV Index Variables Quantile Transform Correlation')
```



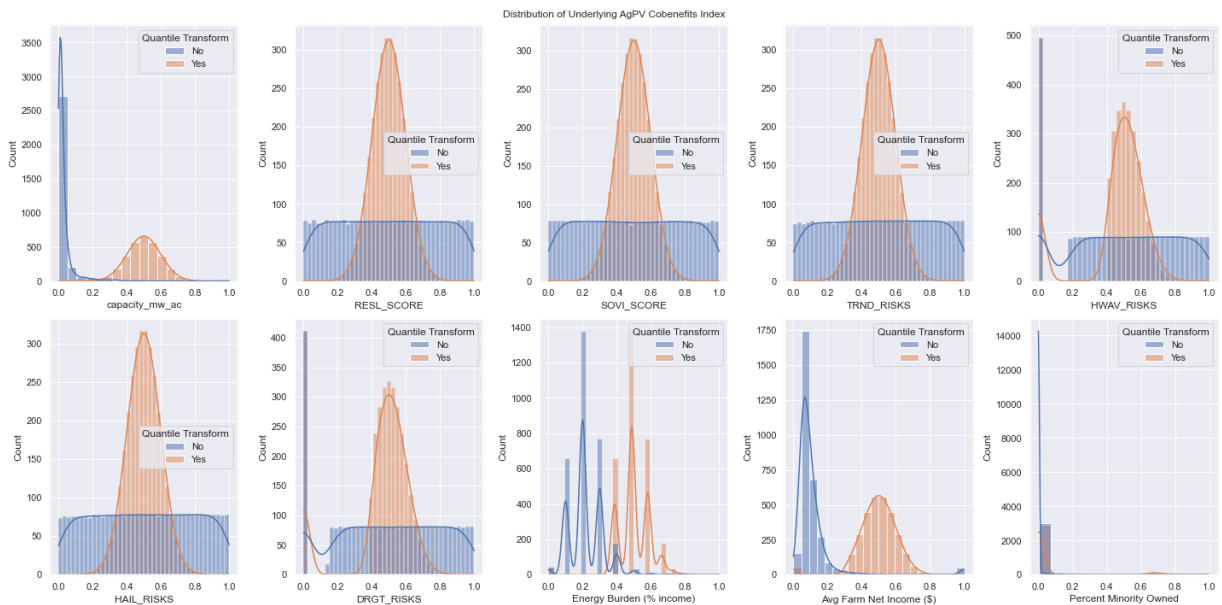
```
In [13]: # Get same columns from normalized df
plot_cols_df = merged_df[normalized_df.columns]
# scale un tranformed data
scaled_original_data = scaler.fit_transform(plot_cols_df)
scaled_original_df = pd.DataFrame(scaled_original_data, index=plot_cols_df.index, c
# add column to signify if this data has been quantile transformed
scaled_original_df['Quantile Transform'] = 'No'
# join dfs
xform_comp_df = pd.concat([scaled_original_df, normalized_df])
xform_comp_df['Quantile Transform'] = xform_comp_df['Quantile Transform'].fillna('Y
```

Look at how wonderful these bell curves look!

```
In [14]: # Look at histograms of transformed data compared to just minmax scaling. Look how
fig, axes = plt.subplots(2, 5, figsize=(20,10))

for col, ax in zip(xform_comp_df.iloc[:, :-1].columns, axes.flatten()):
    sns.histplot(xform_comp_df, x=col, kde=True, ax=ax, hue='Quantile Transform')

fig.suptitle('Distribution of Underlying AgPV Cobenefits Index')
plt.tight_layout()
plt.savefig(os.path.join(graph_dir, 'agpv_variable_distribution.png'))
```



Ag PV Index Calculation

[Apache, AZ](#) stands out due to its near 100% minority cropland ownership, reasonably high energy burden (80th percentile), high drought risk (95th percentile), low torndo risk (2nd percentile), and high solar supply (99th percentile).

Most of this land is owned by the Navajo nation, so the high minoirty ownership is not suprising.

Also note that most of the top counties in overall index seem to have high social vulnerability as well. This makes sense, since climate hazards make up much of the underlying variables.

```
In [15]: # get average of relevant variables
index_vars = ['capacity_mw_ac', 'TRND_RISKS', 'HWAV_RISKS', 'HAIL_RISKS', 'DRGT_RIS',
              'Avg Farm Net Income ($)', 'Percent Minority Owned']

# minmax scale final result
agpv_benefits_score = normalized_df[index_vars].mean(axis=1)
agpv_benefits_score = (agpv_benefits_score - agpv_benefits_score.min()) / (agpv_ben
```

```
In [16]: # merge calculated score with underlying data
score_df = pd.concat([merged_df, agpv_benefits_score], axis=1)
score_df = score_df.rename(columns={0: 'AgPV_cobenefits_score'})
```

```
# filter relevant columns
score_df = score_df[plot_cols + ['AgPV_cobenefits_score']]
score_df = score_df.drop(columns=['capacity_mw_ac'])

# show top n counties
score_df.sort_values(['AgPV_cobenefits_score'], ascending=False).head(10)
```

Out[16]:

	county	state	capacity_mw_dc	RESL_SCORE	SOVI_SCORE	TRND_RISKS	HW
FIPS							
04001	Apache	Arizona	861473.520644	0.640000	96.500000	3.372574	
46071	Jackson	South Dakota	151744.473926	0.380000	97.930000	16.671969	
40135	Sequoyah	Oklahoma	26711.406422	7.100000	86.220001	82.278078	
40107	Okfuskee	Oklahoma	29631.438316	9.930000	86.599998	41.552657	
06107	Tulare	California	38451.738669	12.830000	93.440002	33.439389	
48029	Bexar	Texas	37654.127497	38.639999	92.489998	99.872733	
46121	Todd	South Dakota	119771.505357	0.190000	96.279999	41.170856	
41045	Malheur	Oregon	332408.348564	6.080000	99.430000	10.881324	
51083	Halifax	Virginia	43117.128771	33.669998	70.150002	28.125994	
06047	Merced	California	64317.236668	15.470000	96.980003	28.635062	

In [17]: `score_df.sort_values(['AgPV_cobenefits_score'], ascending=False).rank(pct=True).hea`

Out[17]:


	county	state	capacity_mw_dc	RESL_SCORE	SOVI_SCORE	TRND_RISKS	HWAI
FIPS							
04001	0.022713	0.024335	0.999676	0.006814	0.964958	0.024010	0
46071	0.437216	0.764925	0.957171	0.004218	0.979559	0.156067	0
40135	0.818624	0.691759	0.648605	0.071382	0.863076	0.819598	0
40107	0.674886	0.691759	0.686243	0.098962	0.866321	0.404932	0
06107	0.901687	0.060513	0.773848	0.127515	0.934783	0.322842	0


```
In [18]: fig, ax = plt.subplots(1,1, figsize=(5,10), dpi=300)
g = sns.boxenplot(score_df, y='AgPV_cobenefits_score')

# annotate top n outliers
def annotate_outliers(row, g):
    text = ', '.join([row['county'], row['state']])
    y = row['AgPV_cobenefits_score']
    g.annotate(text, xy=(0.01, y), ha='left', size=7)

score_df.sort_values('AgPV_cobenefits_score', ascending=False).iloc[:3].apply(lambda

fig.suptitle('Ag PV Co-benefits Index Scores')
plt.tight_layout()
plt.savefig(os.path.join(graph_dir, 'agpv_score_boxen.png'))
```





Clustering

Let's do some simple unsupervised learning on this data set. Maybe using clusters can provide more nuanced themes and analysis for groups that a single number might overlook.

This [post](#) has some cool visualization ideas:


```
In [40]: from sklearn.preprocessing import Normalizer, StandardScaler
from sklearn.decomposition import PCA
from sklearn.cluster import KMeans

# normalize data
X_norm = StandardScaler().fit_transform(merged_df[index_vars])

# principal component analysis
pca = PCA()
pca.fit(X_norm)
# get percent variance explained
percent_var = np.round(pca.explained_variance_ratio_ * 100, 3)

sns.lineplot(x=range(1, len(percent_var) + 1), y=percent_var.cumsum(), marker='o',
```

Out[40]: <Axes: >



Without quantile normalization, 6 principle components explain ~85% of the variance. This seems like a good number to choose.

```
In [52]: n_comps=6
pca = PCA(n_components=n_comps)
pca.fit(X_norm)
scores_pca = pca.transform(X_norm)

k_values = range(2,21)
# clustering
WCSS = [] # holds within cluster sum of squares for each value of k

for k in k_values:
    kmeans_pca = KMeans(n_clusters=k, init='k-means++', random_state=123)
    kmeans_pca.fit(scores_pca)
    WCSS.append(kmeans_pca.inertia_)
```

[illegible]

```

C:\Users\stanley27\Anaconda3\lib\site-packages\sklearn\cluster\_kmeans.py:1416: Futu
reWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set t
he value of `n_init` explicitly to suppress the warning
  super()._check_params_vs_input(X, default_n_init=10)
C:\Users\stanley27\Anaconda3\lib\site-packages\sklearn\cluster\_kmeans.py:1416: Futu
reWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set t
he value of `n_init` explicitly to suppress the warning
  super()._check_params_vs_input(X, default_n_init=10)
C:\Users\stanley27\Anaconda3\lib\site-packages\sklearn\cluster\_kmeans.py:1416: Futu
reWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set t
he value of `n_init` explicitly to suppress the warning
  super()._check_params_vs_input(X, default_n_init=10)
C:\Users\stanley27\Anaconda3\lib\site-packages\sklearn\cluster\_kmeans.py:1416: Futu
reWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set t
he value of `n_init` explicitly to suppress the warning
  super()._check_params_vs_input(X, default_n_init=10)
C:\Users\stanley27\Anaconda3\lib\site-packages\sklearn\cluster\_kmeans.py:1416: Futu
reWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set t
he value of `n_init` explicitly to suppress the warning
  super()._check_params_vs_input(X, default_n_init=10)


```

```

In [53]: plt.figure(figsize=(10,6), dpi=300)
g = sns.lineplot(x=k_values, y=WCSS, marker='o', linestyle='--')
g.set_xticks(k_values)
g.set_xlabel('K clusters')
g.set_ylabel('WCSS')

```

```
Out[53]: Text(0, 0.5, 'WCSS')
```



There's no clear elbow point. K=8 or 9 seems like there's a slight change in slope of the line. Using silhouette score might give more detail. Let's use 8 for now.

```

In [54]: k = 8
kmeans_pca = KMeans(n_clusters=k, init='k-means++', random_state=123)

```

```
kmeans_pca.fit(scores_pca)
```

C:\Users\stanley27\Anaconda3\lib\site-packages\sklearn\cluster_kmeans.py:1416: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning
super()._check_params_vs_input(X, default_n_init=10)

Out[54]:

▼ KMeans

KMeans(random_state=123)

In [55]:

```
# Add principal components to original df
pc_col_names = ['PC' + str(k + 1) for k in range(n_comps)]
cluster_df = pd.concat([merged_df[index_vars], pd.DataFrame(scores_pca, columns=pc_col_names)], axis=1)

# add labels from kmeans
cluster_df['Cluster'] = kmeans_pca.labels_
cluster_df['Cluster'] = cluster_df['Cluster'].astype(str)
cluster_df
```

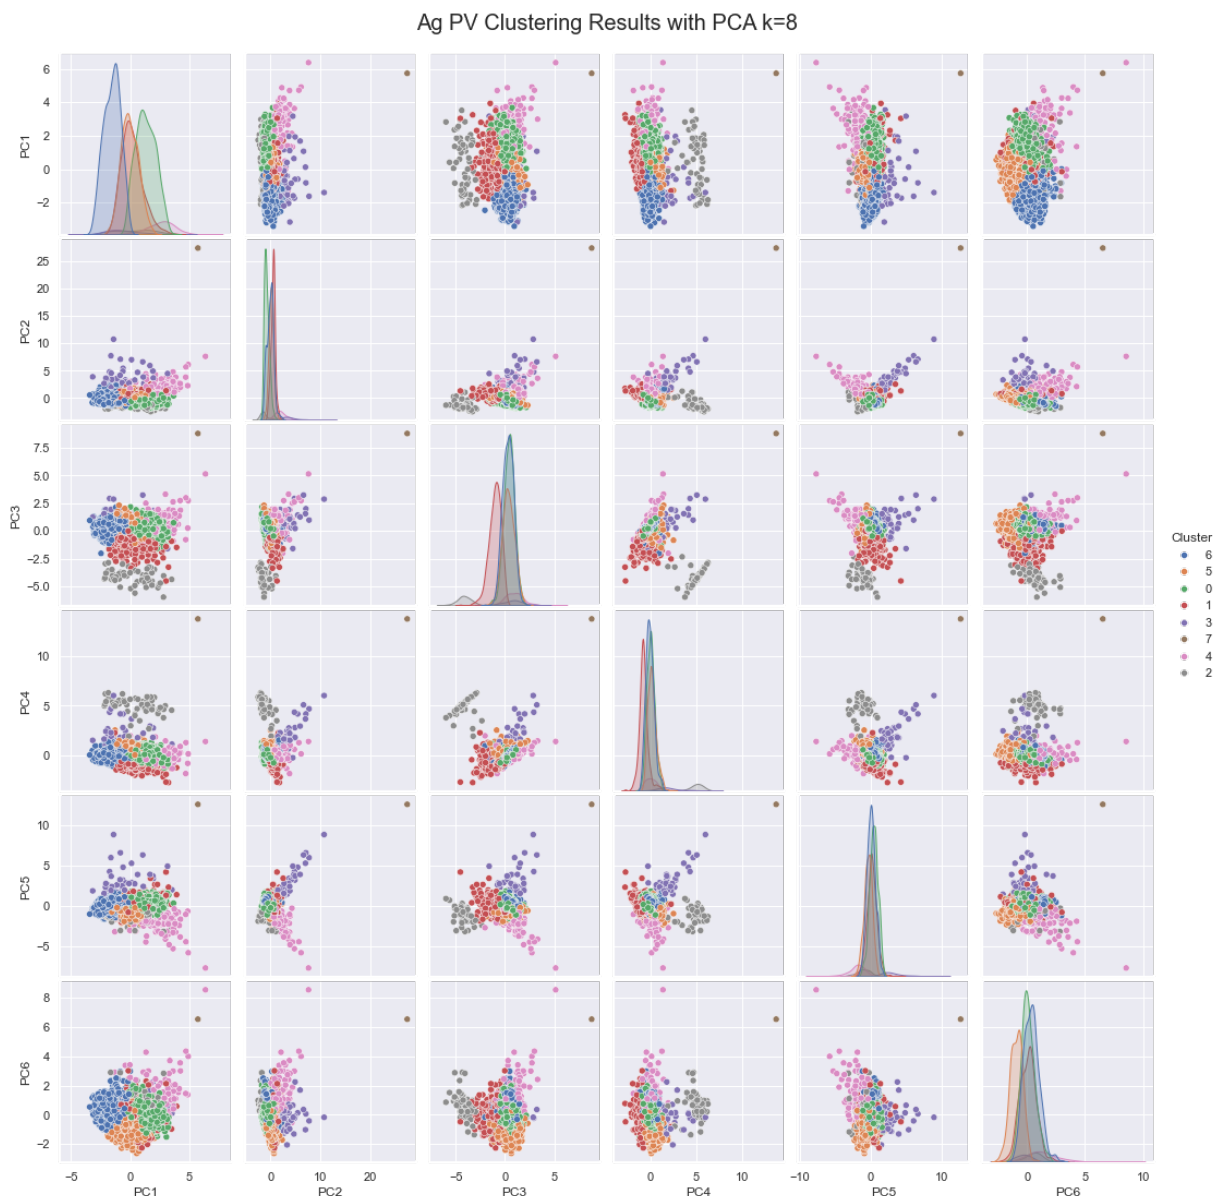
Out[55]:

	capacity_mw_ac	TRND_RISKS	HWAV_RISKS	HAIL_RISKS	DRGT_RISKS	Energy Burden (% income)	Income
FIPS							
01001	22349.202978	73.846643	87.082405	57.111040	50.652243	3.0	18
01003	46948.973155	91.377665	91.091314	52.943048	86.796055	2.0	35
01005	43344.657028	73.942094	69.710468	39.993637	80.687241	4.0	63
01007	28368.690972	71.969456	67.228762	43.684378	32.421254	4.0	-1
01009	18483.385380	89.723194	82.437162	43.716195	52.147630	3.0	66
...
56037	401041.398994	10.531339	0.000000	3.340757	29.812281	2.0	21
56039	4915.661374	21.189946	0.000000	39.102768	21.539930	2.0	6
56041	83816.558073	12.567611	0.000000	17.817372	25.676106	3.0	10
56043	72277.627225	13.013045	0.000000	5.058861	43.143493	4.0	28
56045	155044.608219	24.657970	0.000000	76.073815	19.153675	7.0	15

3082 rows × 15 columns

In [71]:

```
sns.pairplot(cluster_df.iloc[:, -n_comps-1:], hue='Cluster')
plt.suptitle('Ag PV Clustering Results with PCA k=8', size=20, y=1.02)
plt.savefig(os.path.join(graph_dir, 'agpv_pca_clustering.png'))
```



Clustering Analysis

Most of the results seem jumbled together, but there's a few groups that stand out.

- Cluster 2 has clear separation in a few of the views.
- Cluster 4 also gets some separation
- Cluster 7 only has 1 point, but it is a clear outlier from everything else. I bet this is Apache, AZ
- The purple cluster (cluster 3) seems to be closer to cluster 7 than the other clusters

Let's have a look at the clusters:

```
In [67]: cluster_summary = cluster_df.groupby(['Cluster']).mean().round(2)
# drop pc values
cluster_summary = cluster_summary.loc[:, : 'Percent Minority Owned']
cluster_summary
```

Out[67]:

	capacity_mw_ac	TRND_RISKS	HWAV_RISKS	HAIL_RISKS	DRGT_RISKS	Energy Burden (% income)	Ir
Cluster							
0	20948.68	29.27	27.08	30.19	21.01	3.35	
1	23573.55	45.71	36.26	70.49	69.86	4.04	
2	16339.48	44.07	44.14	40.17	24.39	2.75	10
3	28635.92	62.10	72.03	48.65	68.00	3.74	
4	233861.88	14.58	30.13	26.53	61.93	3.31	
5	19520.99	47.86	62.51	29.67	66.85	2.91	
6	16671.52	80.31	70.85	73.52	48.11	2.59	
7	642890.58	3.37	17.82	1.34	95.26	4.00	

Cluster 7 is Apache AZ. High minority cropland ownership!

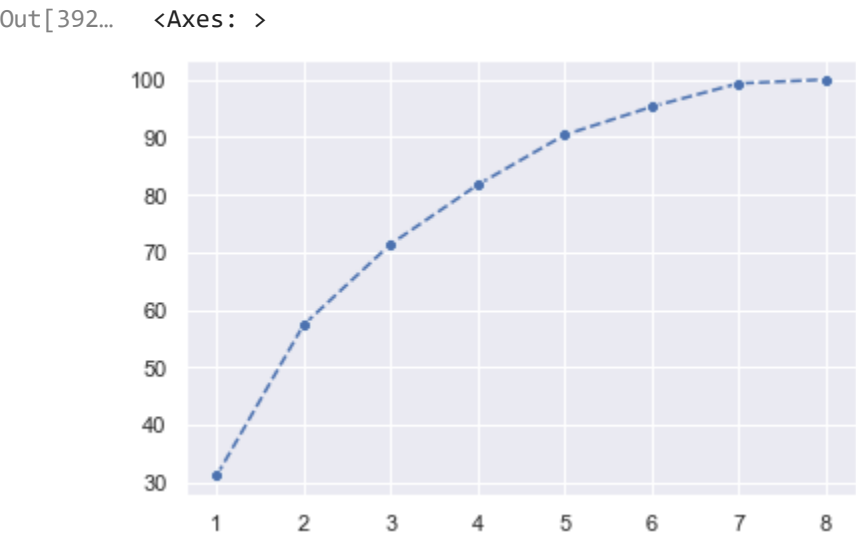
Cluster 4 is defined by it's high solar potential Cluster 2 has high Avg Farm Net income

Cluster 6: high tornado risk

```
In [392... # normalize data
X_norm = Normalizer().fit_transform(normalized_df[index_vars])

# principal component analysis
pca = PCA()
pca.fit(X_norm)
# get percent variance explained
percent_var = np.round(pca.explained_variance_ratio_ * 100, 3)

sns.lineplot(x=range(1, len(percent_var) + 1), y=percent_var.cumsum(), marker='o',
```



In []: